

Multi-Format Document Q&A AI System

Production-Grade Retrieval-Augmented Generation (RAG) Architecture

TARGET ARCHITECTURE	CLOUD BUDGET / SPEND	COMPUTE TARGET	VERIFICATION MODE
RAG (Pinecone + Gemini)	\$0.00 (100% Free Tier)	Zero-GPU (Local CPU)	CLI & Postman REST

1. Executive Summary

This project delivers a complete, enterprise-grade Retrieval-Augmented Generation (RAG) system engineered for high-precision, grounded document question answering. Traditional LLM deployments frequently face critical limitations including hallucinated responses, knowledge cutoff boundaries, and catastrophic context length limitations.

To resolve these constraints at zero infrastructure cost (\$0 cloud spend), we designed and deployed a serverless vector architecture combining **Pinecone Serverless** (cosine distance metric, 768 dimensions) with **Google Gemini API** models (`gemini-embedding-001` and `gemini-3.5-flash`). The system parses heterogeneous document types (`.pdf` and `.json`), constructs semantic chunk vector representations, and executes similarity retrieval with strict guardrailed generation.

768	0.00 s	100%
Vector Dimensions	Local GPU Overhead	Grounded Precision


2. Architectural Justification & Technology Stack

Every technology selection was vetted to achieve strict enterprise criteria: zero ongoing infrastructure cost, enterprise scalability, and deterministic guardrailed outputs.

Layer / Component	Selected Solution	Alternative Considered	Engineering Justification
Vector Database	Pinecone Serverless (AWS us-east-1)	Local ChromaDB / FAISS	Zero persistent RAM footprint on local client; persistent cloud SLA; native managed cosine similarity indexing.
Embedding Engine	Google Gemini (<code>gemini-embedding-001</code>)	Local HuggingFace BERT	Generates 768-dim semantic representations with zero GPU requirement on host; superior domain transfer performance.
Generative LLM	Google Gemini (<code>gemini-3.5-flash</code>)	Ollama (Local Llama-3)	Sub-second token latency, massive context capability, strict adherence to system instructions and temperature=0.0.
Orchestration Environment	Google Antigravity Engine	Manual VS Code setup	Autonomous task execution, built-in dependency handling, direct terminal verification, and reproducible scratch logging.
API Verification	Postman HTTP REST Client	cURL / Browser scripts	Visual header isolation, direct raw JSON body inspection, and independent component validation.

3. System Architecture & End-to-End Pipeline

The system follows a decoupled Two-Phase Pipeline consisting of (1) Ingestion/Indexing and (2) Runtime Retrieval/Generation:



3.1 Chunking & Preprocessing Strategy

Documents are extracted into normalized text streams. For PDFs, `pypdf` reads text blocks across sequential pages. For JSON, structured entities are parsed and formatted into contextual key-value blocks. Text is then chunked via sliding window tokenization (~500 words with 50-word overlap) to guarantee semantic continuity across boundaries.

4. Empirical Verification & Test Case Results

The system was subjected to rigorous validation across in-domain document retrieval and out-of-domain adversarial injection tests.

Test Scenario	Input Query	Top Match & Score	System Response & Verification
In-Domain Grounded (JSON)	"What is the return window and refund policy for items?"	<code>sample_knowledge.json</code> Score: 0.6228	"Based on the provided context, customers can return items within 30 days of purchase for a full refund... Software licenses and digital products are non-refundable once activated. (Source: sample_knowledge.json)"
In-Domain Complex (PDF)	"What is the remote home office setup budget and health insurance coverage?"	<code>company_handbook.pdf</code> Score: 0.7165	"Remote Home Office Setup Budget: Full-time employees are eligible for a one-time remote home office setup budget of \$1,500. Health Insurance: Comprehensive health coverage begins on the first calendar day of month following hire date... (Source: company_handbook.pdf)"
Adversarial / Out-of-Domain	"What is the capital of France?"	<code>N/A (Irrelevant)</code> Score: 0.4880	"Information not found in context" (Hallucination strictly prevented)

5. Postman REST API Technical Specifications

All four core stages of the RAG pipeline were independently validated via direct HTTP REST specifications in Postman.

Request 1: Gemini Vector Embedding Generation

```
POST https://generativelanguage.googleapis.com/v1beta/models/gemini-embedding-001:embedContent?key={{GEMINI_API_KEY}}
Headers: Content-Type: application/json
Body: { "model": "models/gemini-embedding-001", "content": { "parts": [{ "text": "Customers can return items within 30 days of purchase." }] }, "outputDimensionality": 768 }
```

Request 2: Pinecone Vector Record Upsert

```
POST https://rag-knowledge-base-77s36tj.svc.aped-4627-b74a.pinecone.io/vectors/upsert Headers: Api-Key:
{{PINECONE_API_KEY}}, Content-Type: application/json Body: { "vectors": [{ "id": "postman_chunk_001", "values":
[-0.016728, -0.001609, 0.014829, ... 768 float values], "metadata": { "source": "sample_knowledge.json", "text":
"Customers can return items within 30 days of purchase." } } ] }
```

Request 3: Pinecone Vector Similarity Query

```
POST https://rag-knowledge-base-77s36tj.svc.aped-4627-b74a.pinecone.io/query Headers: Api-Key:
{{PINECONE_API_KEY}}, Content-Type: application/json Body: { "vector": [-0.016728, -0.001609, 0.014829, ... 768
float query vector], "topK": 3, "includeMetadata": true, "includeValues": false }
```

Request 4: Grounded Answer Generation (Gemini 3.5 Flash)

```
POST https://generativelanguage.googleapis.com/v1beta/models/gemini-3.5-flash:generateContent?key=
{{GEMINI_API_KEY}} Headers: Content-Type: application/json Body: { "systemInstruction": { "parts": [{ "text": "You
are a strict RAG engine. Answer ONLY using context in <CONTEXT> tags. If missing, output: 'Information not found in
context'." } ] }, "contents": [{ "role": "user", "parts": [{ "text": "<CONTEXT>\n[[SOURCE: sample.json]]\nCustomers
can return items within 30 days.\n</CONTEXT>\n\nUser Question: What is the return window?" } ] } ],
"generationConfig": { "temperature": 0.0 } }
```

6. Operations & Production Roadmap

Operational CLI Execution

To ingest new documentation or query the live index from the terminal:

```
# Ingest new PDF or JSON document python ingest.py data/company_handbook.pdf # Single query execution python
query.py "What is the health insurance coverage policy?" # Interactive chat loop python query.py
```

Future Architectural Roadmap

- **Hybrid Search (Dense + Sparse BM25):** Integrate lexical keyword search with dense embeddings to improve exact-match retrieval for technical part numbers.
- **Cross-Encoder Re-Ranking:** Deploy a lightweight Cohere or BGE re-ranker stage to score the top-20 candidates down to top-3 prior to prompt assembly.
- **Web Interface:** Wrap query engine in a responsive Streamlit or Next.js user dashboard.